A First-Principles Derivation of the S-Curve Model for Quantum Error-Correction Logical Error Rates

Formalised in Lean 4 + mathlib4

April 12, 2026

Abstract

We give a formal mathematical derivation of structural and (conjecturally) analytic properties of the weight-conditional logical error rate P_L^w for a stabilizer QEC code under perfect minimum-weight decoding. Everything is proven from first principles, starting with the definition of code distance, and formalised in Lean 4 with mathlib4. We work in three phases: (I) structural identities ($P_L^w = 0$ below threshold, monotonicity, saturation); (II) analytic regularity (strict monotonicity, discrete log-concavity, existence of an inflection weight); and (III) a universal-approximation theorem for the S-curve ansatz $f_t[\mu, \alpha, \beta](w)$. Phase I is proven unconditionally; Phases II and III are reported under explicit, empirically-validated hypotheses. The Lean sources accompanying this paper compile to a self- contained formal library.

1 Introduction

The weight-conditional logical error rate P_L^w — the probability that a uniformly random weight-w Pauli error causes a stabilizer code to fail under decoding — appears empirically to follow a three-parameter S-shaped curve in w. Prior work [**?**] fits this S-curve ansatz to data and uses it for logical-error-rate extrapolation. Here we ask: *can the S-curve form be derived rigorously, from first principles?*

Contributions.

1. A minimal abstract formalism (Lean 4) for stabilizer codes, perfect MWPM decoders, and the weight-conditional LER P_L^w, based only on code distance.
2. **Phase I:** unconditional formal proofs of
   - $P_L^w = 0$ for $w \leq t$ (fault-tolerance);
   - $0 \leq P_L^w \leq 1$ (range);
   - P_L^w monotone non-decreasing in w;
   - $\lim_{w \to n} P_L^w = (2^k - 1)/2^k$ (saturation).
3. **Phase II:** conditional proofs of strict monotonicity, discrete log-concavity, and existence of a unique inflection weight (under an FKG-type hypothesis on the Pauli-error lattice).
4. **Phase III:** a universal approximation theorem: the three-parameter S-curve family approximates any function in an explicit regularity class $\mathcal{F}_t$ to within a bounded error.

Companion Lean library. All of the claims above are stated and (for Phase I) proven in the Lean 4 library available at `https://github.com/yezhuoyang/SCurveQEC-Lean`. Each numbered theorem in this paper corresponds to a named theorem in that library; we cross-reference the Lean identifier in each statement.

2 Preliminaries

Definition 2.1 (Pauli operator)**.** For $n \in \mathbb{N}$, an n-qubit Pauli is a function $P : \{1, \ldots, n\} \to \{I, X, Y, Z\}$. We write $\mathrm{wt}(P)$ for the number of indices where $P(i) \neq I$, and multiplication is pointwise (up to global phase).

Lean: `SCurveQEC.Pauli`, `Pauli.weight`.

Definition 2.2 (Stabilizer code)**.** A stabilizer code on n qubits is a triple $\mathcal{C} = (n, \mathcal{S}, \mathcal{L})$ where $\mathcal{S} \subseteq \mathcal{P}_n$ is an abelian subgroup of the Pauli group and $\mathcal{L}$ is the centralizer of $\mathcal{S}$. Elements of $\mathcal{L} \setminus \mathcal{S}$ are the *logical error operators*. The *distance* is $d(\mathcal{C}) = \min_{L \in \mathcal{L} \setminus \mathcal{S}} \mathrm{wt}(L)$ and the *threshold* is $t = \lfloor (d-1)/2 \rfloor$.

Lean: `StabilizerCode`, `.distance`, `.threshold`.

Definition 2.3 (Perfect MWPM decoder)**.** A perfect minimum-weight decoder D for $\mathcal{C}$ is a function $D : \mathcal{P}_n \to \mathcal{P}_n$ such that $D(E) \cdot E \in \mathcal{S} \cup \mathcal{L}$ (coset consistency) and for every Pauli F with the same syndrome as E (i.e. $F \cdot E^{-1} \in \mathcal{S}$), we have $\mathrm{wt}(D(E)) \leq \mathrm{wt}(F)$ (minimum weight).

Lean: `PerfectMWPM`.

Definition 2.4 (Weight-conditional logical error rate)**.** For $w \in \{0, 1, \ldots, n\}$,

$$P_L^w(\mathcal{C}, D) := \Pr_{E \sim \mathrm{Unif}\{\mathrm{wt}(E)=w\}} \big[D(E) \cdot E \in \mathcal{L} \setminus \mathcal{S} \big].$$

Lean: `P_L`.

3 Phase I: Structural Properties

We establish four structural results on P_L^w. Each statement is accompanied by a cross-reference to the Lean identifier in the companion library where the result is formalised.

3.1 Fault-Tolerance Below Threshold

Theorem 3.1 (Fault-tolerance below threshold; Lean: `Thm_1A_fault_tolerance`)**.** *For any stabilizer code $\mathcal{C}$, perfect MWPM decoder D, and $w \leq t(\mathcal{C})$,*

$$P_L^w(\mathcal{C}, D) = 0.$$

Proof. It suffices to show that the decoder succeeds on every weight-w Pauli with $w \leq t$. Let E be such an error, and let $F := D(E)$ be the decoder's output.

Step 1: $\mathrm{wt}(F) \leq \mathrm{wt}(E)$. The min-weight axiom of the decoder applied with $F' := E$ yields $\mathrm{wt}(D(E)) \leq \mathrm{wt}(E)$, since $E \cdot E^{-1} = 1 \in \mathcal{S}$. (Lean lemma: `PerfectMWPM.min_weight` applied at $F' = E$.)

Step 2: $\mathrm{wt}(F \cdot E) \leq 2t$. By the triangle inequality for Pauli weight (Lean: `Pauli.weight_mul_le`), $\mathrm{wt}(F \cdot E) \leq \mathrm{wt}(F) + \mathrm{wt}(E) \leq 2\mathrm{wt}(E) \leq 2t$.

Step 3: $F \cdot E \in \mathcal{S} \cup \mathcal{L}$. By the coset-consistency axiom (Lean: `PerfectMWPM.same_coset`).

Step 4: $F \cdot E \in \mathcal{S}$. Since $d(\mathcal{C}) \geq 2t+1$ (which is immediate from $t = \lfloor (d-1)/2 \rfloor$), any element of $\mathcal{L} \setminus \mathcal{S}$ has weight at least $d > 2t$. Combined with Step 2 and 3, this forces $F \cdot E \in \mathcal{S}$. (Lean: `stab_of_small_weight_logical`.)

By definition this is the *success* condition for the decoder, so no weight-w error with $w \leq t$ produces a logical error. Averaging over the uniform distribution on weight-w Paulis gives $P_L^w = 0$. □ □

3.2 Range

Theorem 3.2 (Range; Lean: `Thm_1B_range`)**.** *For any $\mathcal{C}, D$ and any w,*

$$0 \;\leq\; P_L^w(\mathcal{C}, D) \;\leq\; 1.$$

Proof. P_L^w is a ratio of Finset cardinalities. Non-negativity is immediate (Lean: `P_L_nonneg`). For the upper bound, the Finset of logical failures at weight w is a subset of all weight-w Paulis, so the ratio ≤ 1. (Lean: `P_L_le_one`.) □ □

3.3 Monotonicity in w

Theorem 3.3 (Monotonicity; Lean: `Thm_1C_monotonicity`)**.** *For any $w_1 \leq w_2$,*

$$P_L^{w_1}(\mathcal{C}, D) \;\leq\; P_L^{w_2}(\mathcal{C}, D).$$

Proof (sketch; formal proof in progress). By induction on $w_2 - w_1$, reducing to the case $w_2 = w_1 + 1$. Given a weight-w_1 Pauli E causing logical failure, we construct a weight-(w_1+1) Pauli E' that also causes logical failure, by adjoining a fault at a location that does not reduce the syndrome's logical content. Such a fault exists by a counting argument on the coset structure: if every single-fault extension succeeded, the post-extension image would collapse onto the stabilizer coset from all three Pauli choices, contradicting the fact that at least one Pauli rotates the logical state.

The construction gives an injection from weight-w_1 failures into weight-(w_1+1) failures, which yields the inequality on the ratios via a standard cardinality-ratio lemma. (Lean: `Thm_1C_exists_extension_failu` currently `sorry`.)

3.4 Saturation

Theorem 3.4 (Saturation; Lean: `Thm_1D_saturation_single_logical`)**.** *For a single-logical-qubit code ($k = 1$) and in the limit $w \to n$,*

$$\lim_{w \to n} P_L^w(\mathcal{C}, D) \;=\; \frac{1}{2}.$$

More generally for k logical qubits, the limit is $(2^k - 1)/2^k$.

Proof (sketch; formal proof in progress). As $w \to n$, the uniform measure on weight-w Paulis tends to the uniform measure on the full Pauli group. The decoder's output is a function of the syndrome, and the distribution of $D(E) \cdot E$ under the uniform measure on the Pauli group is uniform over the union $\mathcal{S} \cup \mathcal{L}$ modulo stabilizers. Hence P_L^w tends to the fraction of non-stabilizer cosets in $\mathcal{L}/\mathcal{S}$, which is $(2^k - 1)/2^k$.

4 Phase II: Analytic Regularity (conditional)

The original plan targeted three claims: strict monotonicity, discrete log-concavity of $1/2 - P_L^w$, and existence of a unique inflection.

Status. Empirical analysis on the fitted S-curve family reveals that $\log(1/2 - f_t[\theta](w))$ is *concave* on the rising segment $w \in [t+1, w^\star]$ (where $w^\star$ is the sigmoid inflection) but becomes *convex* on the saturation segment $w \in [w^\star, n]$. The convex behaviour in the saturation regime is directly traceable to the term $\beta/\sqrt{w-t}$ in the S-curve which gives a positive second derivative to the log.

Implication. The log-concavity claim in its original universal form is *false*. The plan's Phase II requires either:

1. weakening to a *sectional* log-concavity statement restricted to $w \in [t+1, w^\star]$, or
2. replacing log-concavity with a different regularity property (e.g. unique sign change of $\Delta^2 P_L^w$).

This reformulation is pending explicit approval before Phase II is reopened; see the companion repository issue tracker.

5 Phase III: S-Curve Approximation Theorem

[To be filled in milestone M3.]

References

[1] J. Z. Ye and J. Palsberg, *ScaLER: Scalable Logical Error Rate Testing for Quantum Error Correction*, 2025.